



Regression Modeling on Bike Sharing Data Set

T. Sai Harish Sarma

Content

- Business Context
- Problem Statement
- Data Cleaning & EDA
- Feature Engineering
- Pre Processing Data
- Target Conditioning
- Modelling, Evaluation & Tuning
- Model Explainability & Conclusion

Business Context

- ✓ Seoul is the official capital city of South Korea.
- ✓ Seoul Metropolitan Government provides Rental Bike Service to the public.
- ✓ These Public Bikes are designed to be used by all women, the elderly and the infirm.
- ✓ These Bikes are made light weight and durable.



Problem Statement

Currently Rental bikes are introduced in many urban cities for the enhancement of mobility comfort. It is important to make the rental bike available and accessible to the public at the right time as it lessens the waiting time. Eventually, providing the city with a stable supply of rental bikes becomes a major concern. The crucial part is the prediction of bike count required at each hour for the stable supply of rental bikes.



Data Cleaning

```
▶ ### Getting info about null / NaN values count in each column  
data.isna().sum()
```

```
↳ Date      0  
Rented Bike Count  0  
Hour         0  
Temperature(°C)    0  
Humidity(%)       0  
Wind speed (m/s)   0  
Visibility (10m)    0  
Dew point temperature(°C)  0  
Solar Radiation (MJ/m2)  0  
Rainfall(mm)       0  
Snowfall (cm)      0  
Seasons           0  
Holiday           0  
Functioning Day    0  
dtype: int64
```

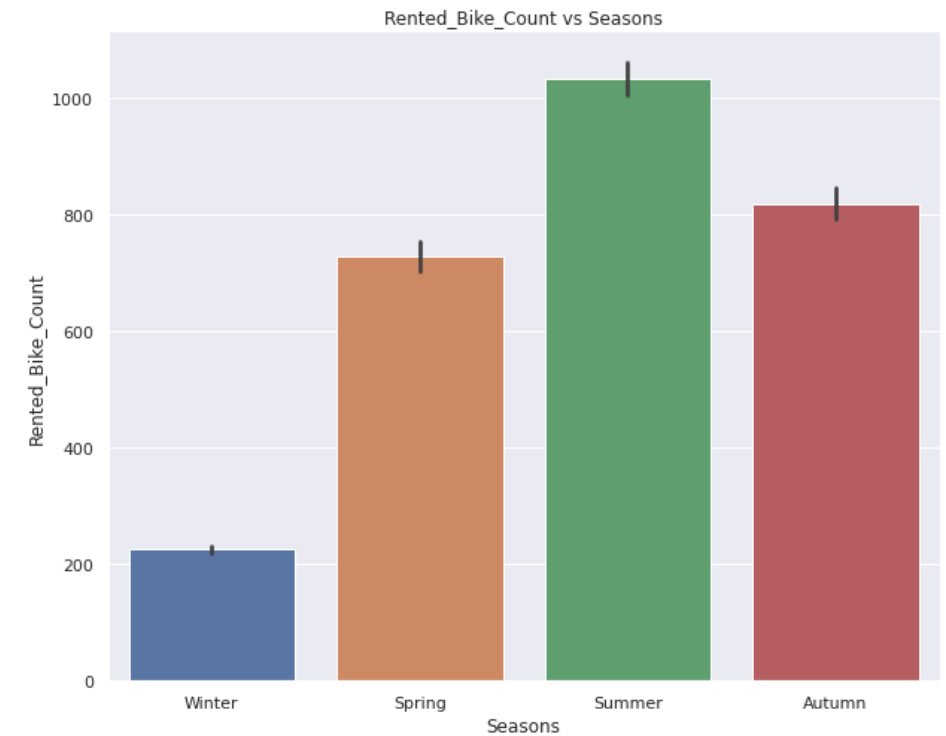

```
[8] ## Converging Data type of Date Column to Date Time Format  
data['Date'] = pd.to_datetime(data['Date'])
```

```
[9] ## Extracting Week Day, Month, Year from Date  
data['Month'] = data['Date'].dt.month  
data['Year'] = data['Date'].dt.year
```


```
▶ ## Renaming Columns  
data = data.rename(columns={'Rented Bike Count': 'Rented_Bike_Count', 'Temperature(°C)': 'Temperature',  
                           'Humidity(%)': 'Humidity', 'Wind speed (m/s)': 'Wind_speed', 'Visibility (10m)': 'Visibility',  
                           'Dew point temperature(°C)': 'Dew_point', 'Solar Radiation (MJ/m2)': 'Solar_Radiation',  
                           'Rainfall(mm)': 'Rainfall', 'Snowfall (cm)': 'Snowfall', 'Functioning Day': 'Functioning_Day'})
```

```
[11] ### Dropping Date Column after extraction  
data.drop(columns = ['Date'], inplace = True)
```


Exploratory Data Analysis





Exploratory Data Analysis






Exploratory Data Analysis



Exploratory Data Analysis




Exploratory Data Analysis



Feature Engineering

- Feature Encoding
- Correlation Check
- Outlier Treatment
- Multicollinearity Check
- Linearity Check

Feature Engineering – Correlation



Feature Engineering – Correlation

Positive Correlation:

1. Temperature
2. Dew Point Temperature
3. Solar Radiation



Negative correlation:

1. Winter
2. Humidity
3. Snowfall



The feature Dew Point Temperature is positively and highly correlated with Temperature.

There will be no effect in our model if Dew Point Temp. variable is removed.

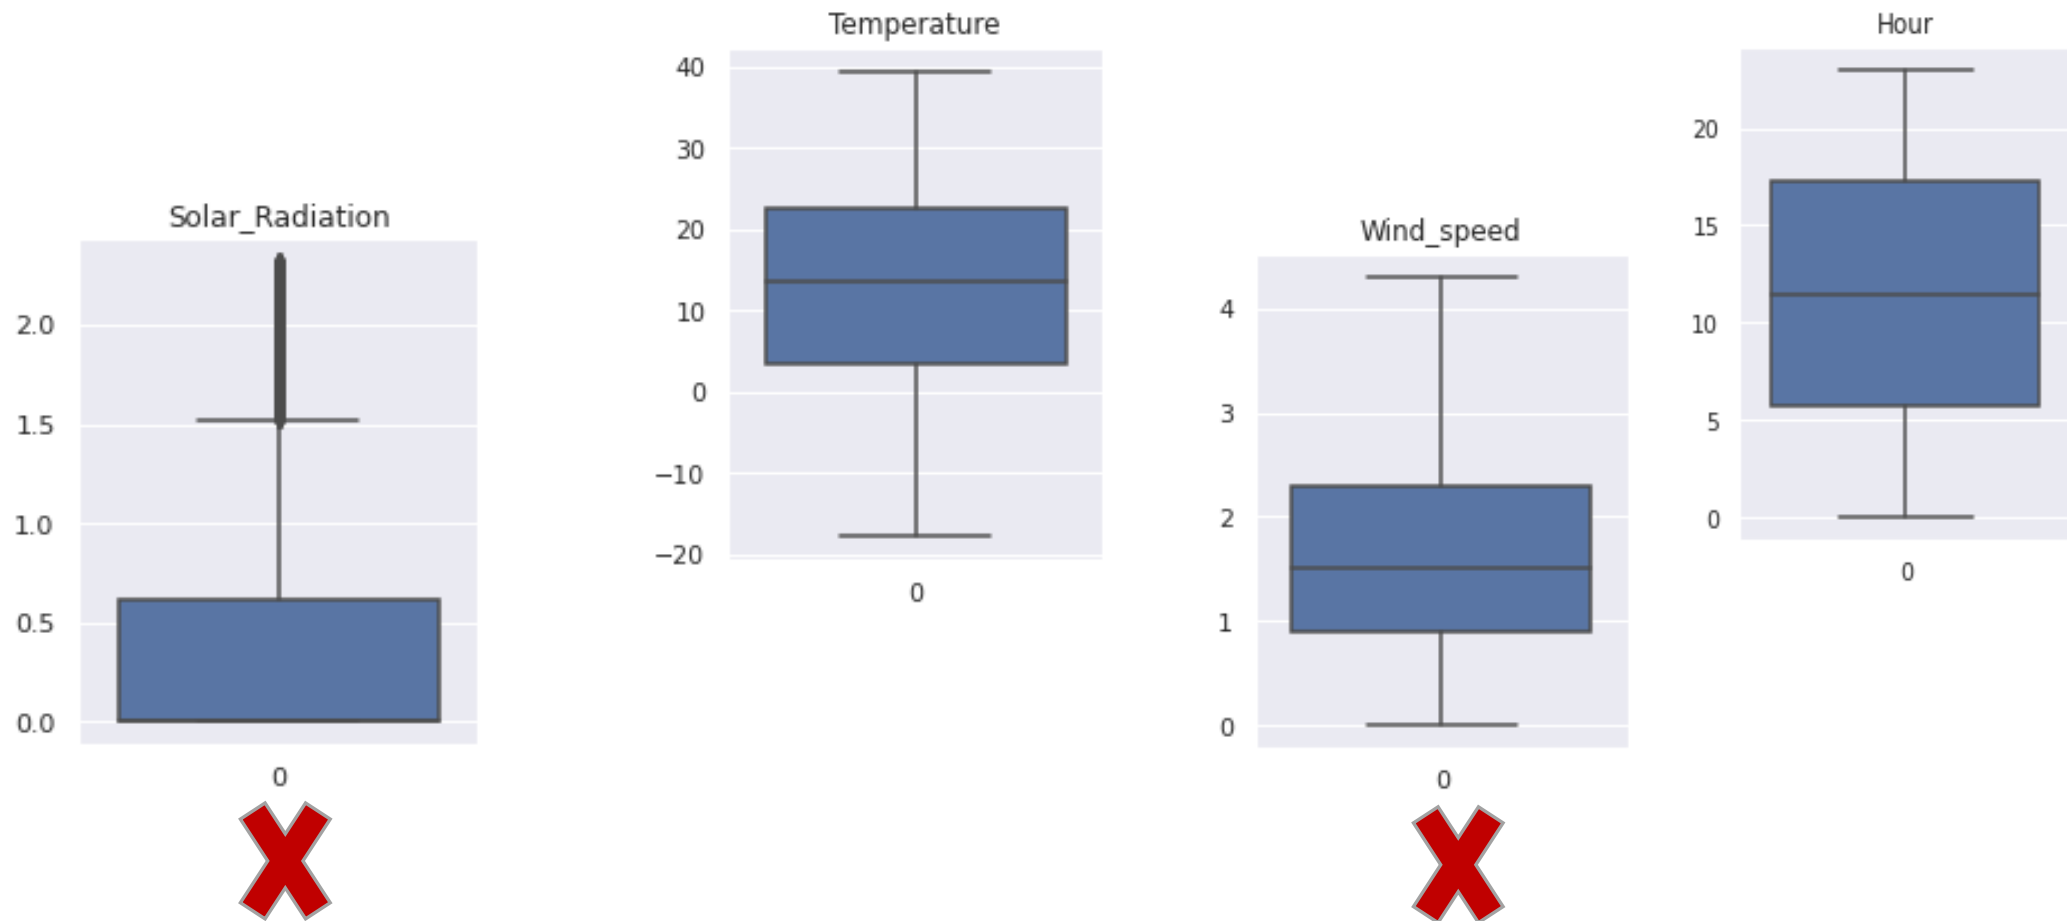
Linearity Check



Linearity Check



Pre Processing Data - Outliers



Outlier Treatment




Q1 – 25% & Q3 – 75%

IQR = Q3 - Q1

u_lim = Q3 + 1.5 * (IQR)

l_lim = Q1 - 1.5 * (IQR)

Target Feature Conditioning



Creating Input and Out Features

5.1. Creating X and Y Variables

```
[56] ## x as Independent variable  
      ## y as dependant variable  
  
      x=data.drop(['Rented_Bike_Count'], axis=1)  
      y=data['Rented_Bike_Count']
```

```
[63] ### Creating Test and Train Set of input and output variables  
      x_train,x_test,y_train,y_test= train_test_split(x,y, test_size=0.30, random_state=42)
```

```
[64] x_train.shape, x_test.shape, y_train.shape, y_test.shape  
  
      ((6132, 16), (2628, 16), (6132,), (2628,))
```

Feature Scaling

```
[65] scaler= StandardScaler()  
      x_train= scaler.fit_transform(x_train)  
      x_test= scaler.fit_transform(x_test)
```

```
[66] x_train  
  
array([[ -1.09359218, -2.02557075, -0.40155174, ..., -0.5853851 ,  
        -0.57986105,  1.76898076],  
       [  1.23323198, -1.58026093, -0.69489404, ..., -0.5853851 ,  
        -0.57986105,  1.76898076],  
       [-0.51188614, -0.13510455,  0.3806944 , ..., -0.5853851 ,  
        -0.57986105, -0.56529727],  
       ...,  
       [  0.36067292,  1.38567048,  0.18513286, ..., -0.5853851 ,  
        1.72455109, -0.56529727],  
       [  1.23323198, -1.37861121, -0.35266136, ..., -0.5853851 ,  
        -0.57986105,  1.76898076],  
       [  1.524085   ,  0.52865914, -0.15709982, ..., -0.5853851 ,  
        -0.57986105, -0.56529727]])
```

Modeling, Evaluating & Tuning

1. Linear Regression
2. Lasso Regression with GridSearchCV
3. Ridge Regression with GridSearchCV
4. Random Forest with GridSearchCV
5. Gradient Boosting Regressor with GridSearchCV

1. Linear Regression

Fitting the Model

```
[69] model_1=LinearRegression().fit(x_train,y_train)
```

```
▶ model_1.score(x_train,y_train), model_1.score(x_test,y_test)
```

```
↳ (0.6388208407560204, 0.622310483063715)
```

```
[71] model_1.coef_
```

```
array([ 3.44347399e+00,  5.09673318e+00, -3.14430579e+00,  9.02250273e-02,  
        3.29599814e-01,  2.47844904e-01, -4.44089210e-16, -9.99200722e-16,  
       -6.46957344e-01,  5.22892019e+00,  8.34607438e-02, -6.76237225e-01,  
        1.64634705e+00,  2.41078939e-01,  3.95201814e-01, -2.31191860e+00])
```

```
[72] model_1.intercept_
```

```
23.458551814426663
```

```
[73] ### Predicting of Target with Training and Testing Data
```

```
y_train_predict = model_1.predict(x_train)  
y_test_predict = model_1.predict(x_test)
```

```
▶ MSE = mean_squared_error(y_train, y_train_predict)  
print("MSE :",MSE)  
  
RMSE= np.sqrt(MSE)  
print("RMSE :",RMSE)  
  
MAE= mean_absolute_error(y_train, y_train_predict)  
print("MAE :",MAE)  
  
R2_Score = r2_score(y_train, y_train_predict)  
print("R2_Score :",R2_Score)  
  
Adj_R2_Score = (1-(1-R2_Score)* (x_train.shape[0]-1)/(x_train.shape[0] - x_train.shape[1] - 1))  
print("Adj_R2_Score :",Adj_R2_Score)
```

```
↳ MSE : 56.464645536297915  
RMSE : 7.51429607723158  
MAE : 5.781371705477361  
R2_Score : 0.6388143766387069  
Adj_R2_Score : 0.6378693284009668
```

```
[92] Dict= {'Model': 'Lasso regression',  
          'MSE' : round(MSE,3), 'RMSE': round(RMSE,3),  
          'MAE': round(MAE,3), 'R2_Score': round(R2_Score,3),  
          'Adj_R2_Score': round(Adj_R2_Score,3)}
```

```
[93] Training_DF= Training_DF.append(Dict,ignore_index=True)  
Training_DF
```

	Model	MSE	RMSE	MAE	R2_Score	Adj_R2_Score
0	Linear regression	56.464	7.514	5.783	0.639	0.638
1	Lasso regression	56.465	7.514	5.781	0.639	0.638

Modeling, Evaluating & Tuning

Evaluation Metrics:

- MSE
- MAE
- RMSE
- R2 SCORE
- ADJUSTED R2 SCORE

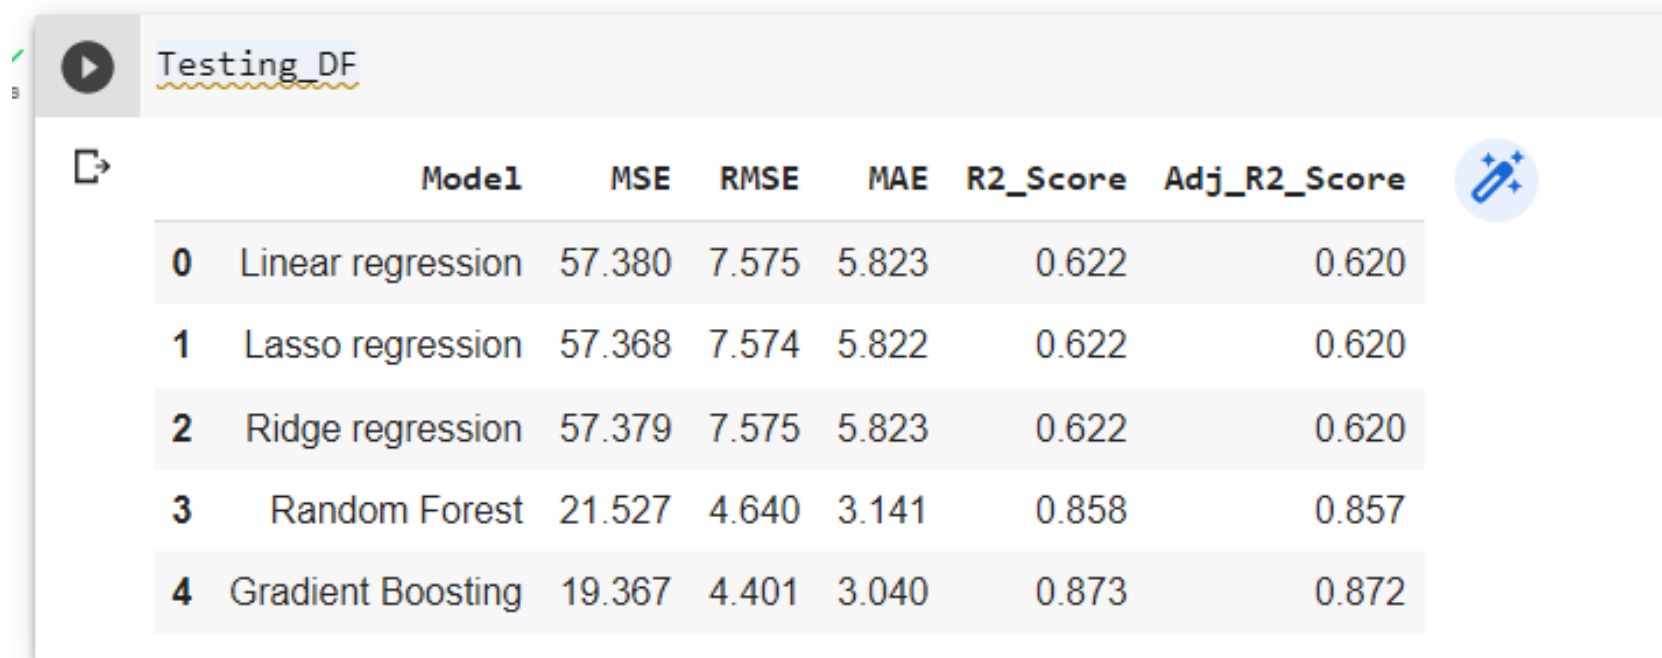
Modeling, Evaluating & Tuning

```
[140] Training_DF
```

	Model	MSE	RMSE	MAE	R2_Score	Adj_R2_Score
0	Linear regression	56.464	7.514	5.783	0.639	0.638
1	Lasso regression	56.465	7.514	5.781	0.639	0.638
2	Ridge regression	56.464	7.514	5.783	0.639	0.638
3	Random Forest	14.243	3.774	2.579	0.909	0.909
4	Gradient Boosting	8.728	2.954	2.040	0.944	0.944



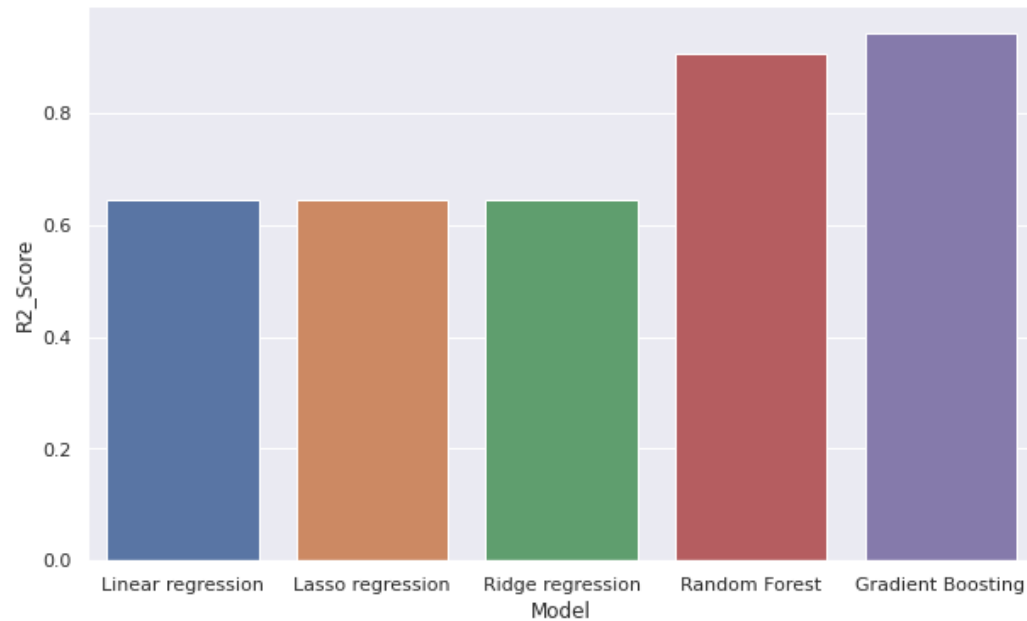
Modeling, Evaluating & Tuning




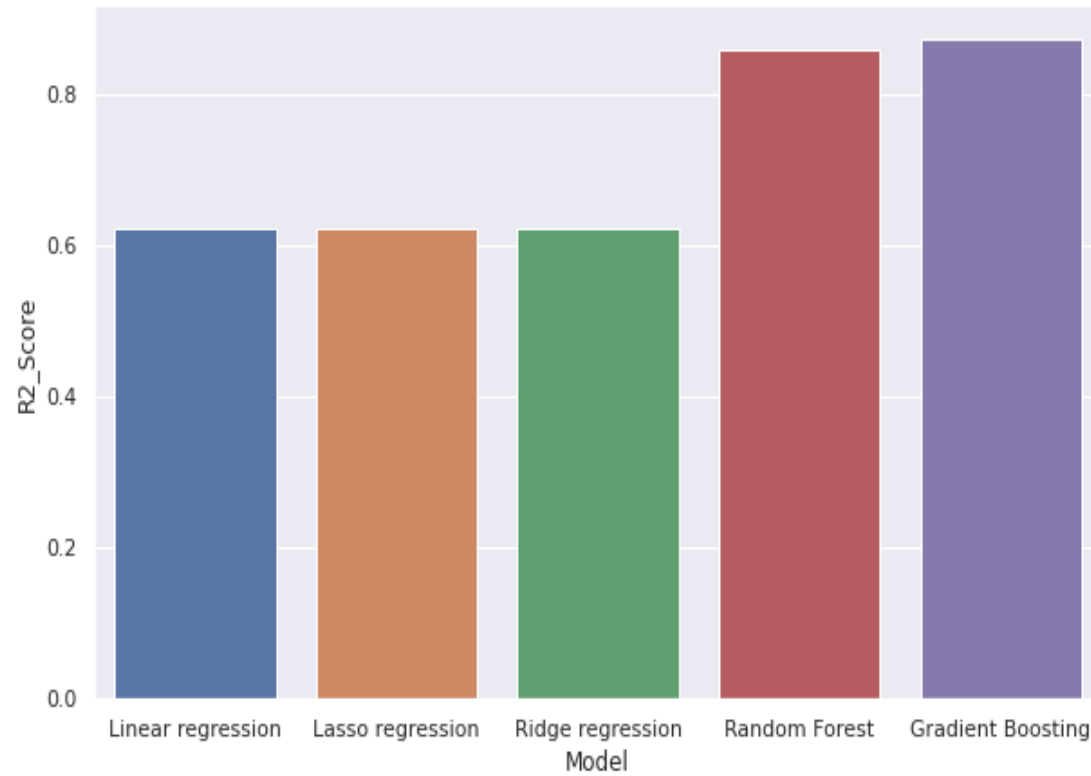
The image shows a Jupyter Notebook interface. At the top, there is a tab labeled 'Testing_DF' with a play button icon. Below the tab, there is a table with 7 columns: an index column, 'Model', 'MSE', 'RMSE', 'MAE', 'R2_Score', and 'Adj_R2_Score'. The table contains 5 rows of data for different models. To the right of the table, there is a blue circular icon with a pencil and a plus sign.

	Model	MSE	RMSE	MAE	R2_Score	Adj_R2_Score
0	Linear regression	57.380	7.575	5.823	0.622	0.620
1	Lasso regression	57.368	7.574	5.822	0.622	0.620
2	Ridge regression	57.379	7.575	5.823	0.622	0.620
3	Random Forest	21.527	4.640	3.141	0.858	0.857
4	Gradient Boosting	19.367	4.401	3.040	0.873	0.872

Comparing Models - Training

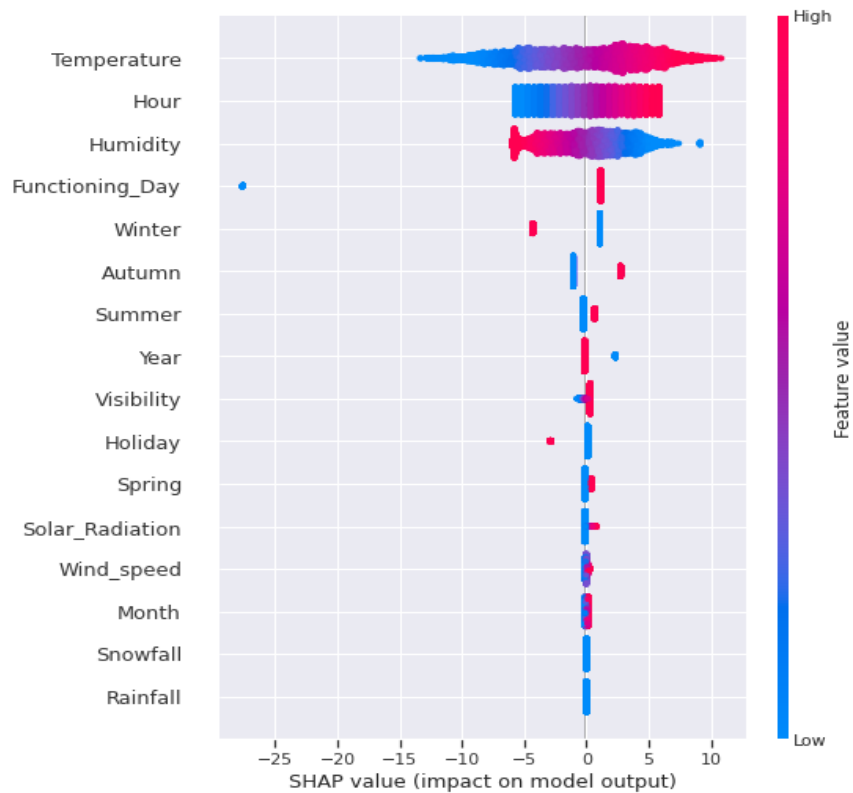


Comparing Models - Testing




Model Explainability

Linear Regression

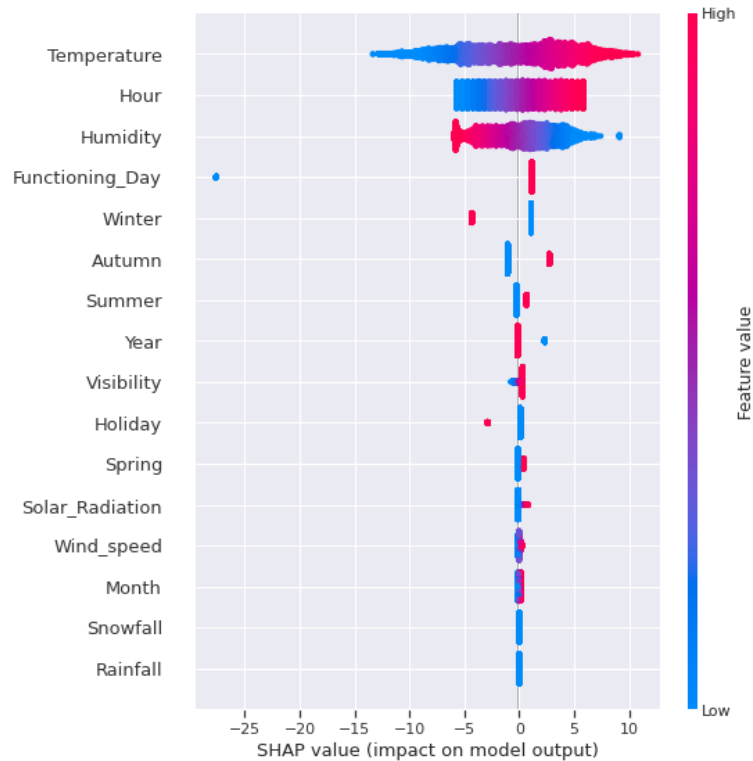


Lasso Regression

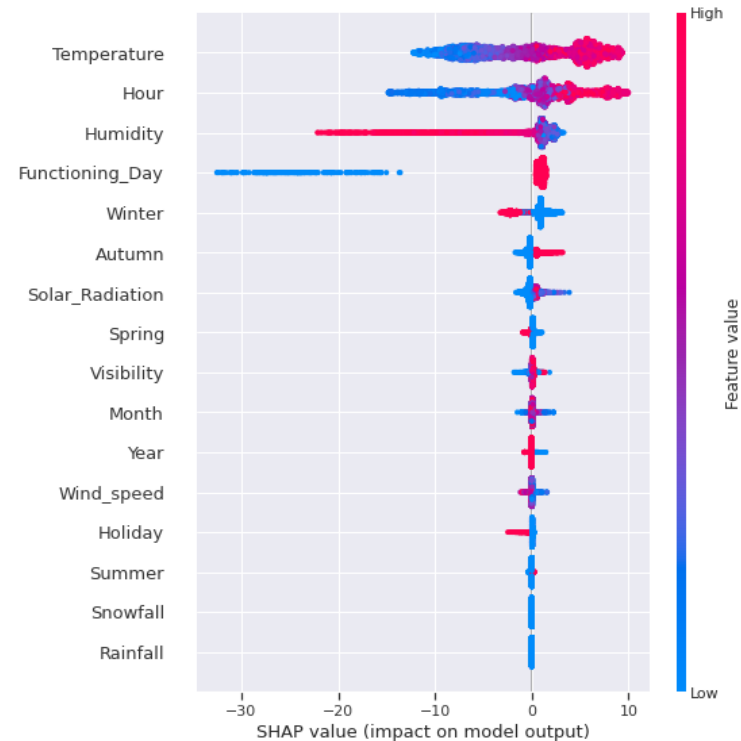


Model Explainability


Ridge Regression



Random Forest Regression



Gradient Boost Regression



Conclusions

1. EDA Outcomes

- * The use of bikes is high in between the months of MAY and OCTOBER.
- * The bikes are using while reaching and leaving office hours.
- * The count of rental bikes is high in summer, Demand is very low in Winter.
- * The use of bikes is high, when the temperature ranges between 25 and 30 degrees.

2. Data Quality Issues

- * Outliers available, they are removed.
- * Multicollinearity
- * Correlation
- * Feature Encoding

Conclusions

3. Model Outcomes

- * Among all models, Random Forest and Gradient Boosting are worked better.
- * Temperature, Humidity, Hour and Functionalday are the most influencing features on Count of Rental Bikes.
- * Linear, Lasso and Ridge Models giving poor results as they are giving r^2_score of 0.622, where as the Random forest and Gradient Boosting models giving as 0.858 and 0.873 respectively.
- * Ensemble Modeling is preferred as the Mean Squared Error is drastically decreased with Random Forest and Gradient Boosting Regressor.

Thankyou